


# Deep Learning Project I

Email:

[[gchalkiopoulos@aueb.gr](mailto:gchalkiopoulos@aueb.gr) - p3352124, [ntsoukalelis@aueb.gr](mailto:ntsoukalelis@aueb.gr) - p3352123]

## MLPs and CNN for image classification on fashion-mnist dataset



Team members: Georgios Chalkiopoulos, Nikolaos  
Tsoukalelis

# CONTENTS

Problem Description ..... 2

Weights & Biases .....3

Visualize .....4

Multilayer Perceptron (MLP) .....5

Convolutional Neural Network (CNN) ..... 10


# Problem Description

## Assignment and Dataset description

Given an image of a fashion item, we had to build a deep learning model that recognizes the fashion item. We had to use at least 2 different architectures, one with MLPs and one with CNNs. We used the Fashion-MNIST dataset to train and evaluate our models. More information about the task and the dataset can be found at <https://github.com/zalando-research/fashion-mnist>. The dataset is also available from Tensorflow.

Fashion-MNIST is a dataset of Zalando's article images—consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes. It shares the same image size and structure of training and testing splits.

Here's an example of how the data looks (each class takes three-rows):



# Weights & Biases

## Weights & Biases API description

In order to make experiment tracking easier we will use [Weights & Biases](https://wandb.ai/aueb), which offers a free licence for academic purposes. For the sake of this assignment a team has been created <https://wandb.ai/aueb>. Access can be granted by contacting the authors.

Note that this code assumes that you have already set up a Wandb account and API key. If you haven't done so yet, you will need to sign up for a free account at <https://wandb.ai/> and follow the instructions there to obtain your API key.

In [4]:

```
wandb.login()
```

```
wandb: Logging into wandb.ai. (Learn how to  
deploy a W&B server locally: https://wandb.  
me/wandb-server)
```

```
wandb: You can find your API key in your br  
owser here: https://wandb.ai/authorize
```

```
wandb: Paste an API key from your profile a  
nd hit enter, or press ctrl+c to quit:
```

```
wandb: Appending key for api.wandb.ai to yo  
ur netrc file: /root/.netrc
```

Out[4]:

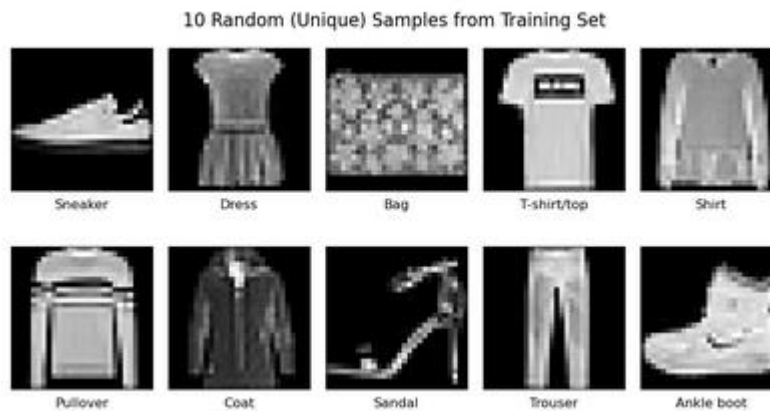
```
True
```

# Visualize

## Training Dataset

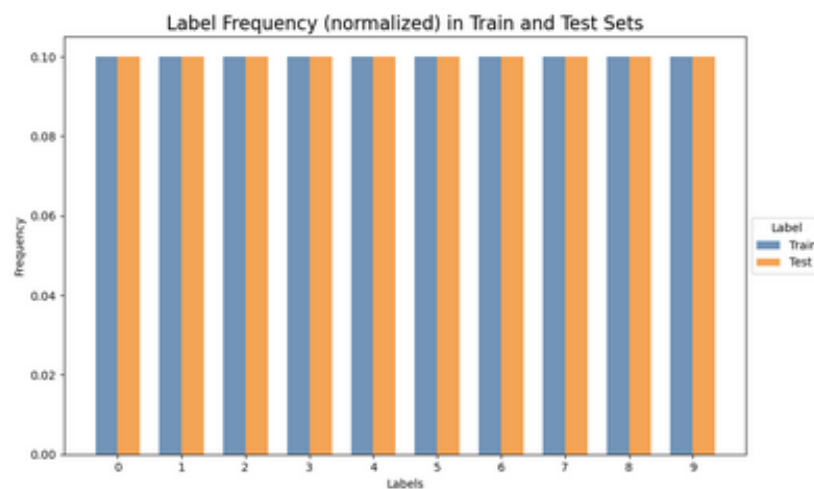
Before proceeding in the implementation of the model and further work on the assignment we produced some visualizations regarding the dataset.

Firstly we produce random samples from Training Set.



## Imbalance

Then we also checked the imbalance of our data and produced the labels' frequency.



# Multilayer Perceptron (MLP)

5

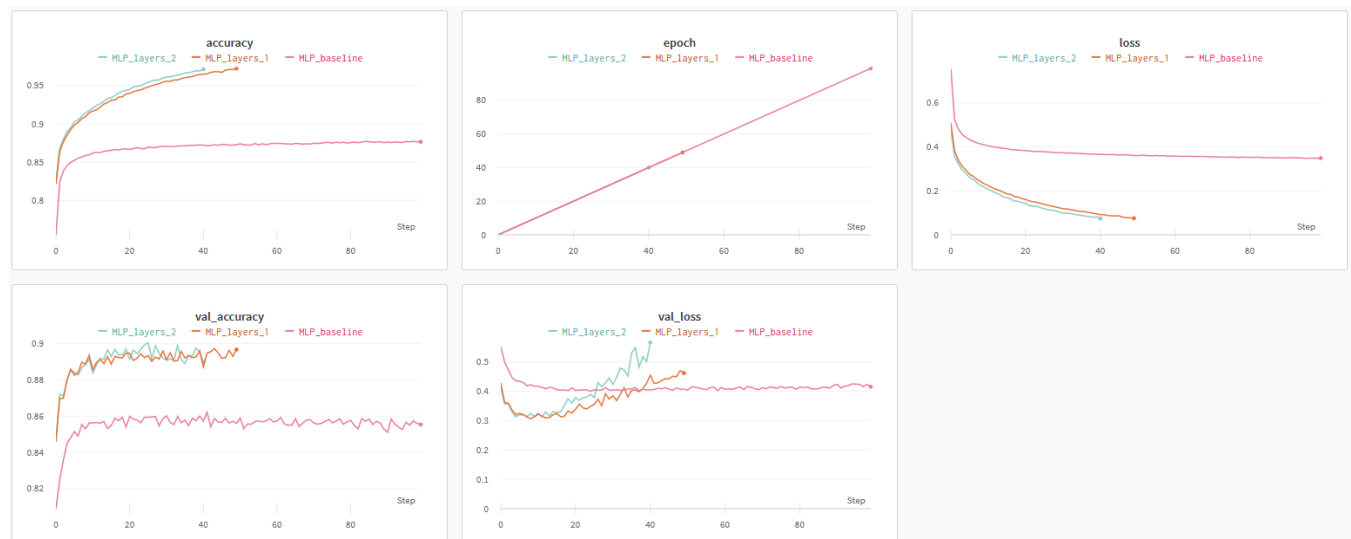
## Baseline MLP model

In this part of the assignment we will evaluate the performance of an MLP architecture. We, firstly, built a Logistic Regression model, which served as a baseline for our process.

Furthermore with the use of Weights and Biases we managed to visualize the progress for any model implementation we applied. The implementation of our baseline model:



We configured the base model by adding 1 (with dimensions [256]) and 2 (with dimensions [512] [256]) hidden layers respectively to see its behavior. The results in comparison with the base model:



# Multilayer Perceptron (MLP)

## Model optimization


It looks that by adding a hidden layer the model suffers from overfitting. This can be seen by comparing the training loss (loss in Weights and Biases) to the validation loss (val\_loss in Weights and Biases). Moreover, it seems that there isn't much of an advantage when adding two hidden layers, compared to one.

We tried to fix these issues, and improve the performance, by performing various optimizations, like:

- Using regularization techniques
- Introduce early stopping
- Add dropout and/or batch normalization
- Add data augmentation
- Reduce the size of the hidden layer
- Optimize hyper-parameters

Following on we applied some sweeps using the Weights and Biases. Actually our approach was to test different combinations of configurations that can be found in our python code, and by showcasing the results in Weights and Biases, to find and highlight the best model. Then we loaded the parameters of this model in our code and used it as final model.

Below we can see, after the aforementioned implementations and having performed a large number of runs, we will stop further exploration, since the results don't seem to improve dramatically. the best configuration and its respective parameters that we eventually used for our MLP model:



As can be seen, above, the best configuration is **youthful-sweep-13**.

# Multilayer Perceptron (MLP)

7

We then loaded this best model-configuration and evaluated the results of the model on the test data:

```
wandb: 5 of 5 files downloaded.
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
Input (InputLayer)	[(None, 784)]	0
Hidden-1 (Dense)	(None, 1024)	803840
Dropout-1 (Dropout)	(None, 1024)	0
Output (Dense)	(None, 10)	10250

```
=====
```

```
Total params: 814,090
```

```
Trainable params: 814,090
```

```
Non-trainable params: 0
```

```
-----
```

```
None
```

```
313/313 [=====] - 4s 3ms/step - loss: 0.2921 - accuracy: 0.9008
```

```
---
```

```
Test Loss      : 0.29207
```

```
---
```

```
Test Accuracy   : 0.90080
```

We have reached to a satisfactory score after our optimization process. Next step was to evaluate the results per item as well. From the images provided, we expect some items to be difficult to distinguish.

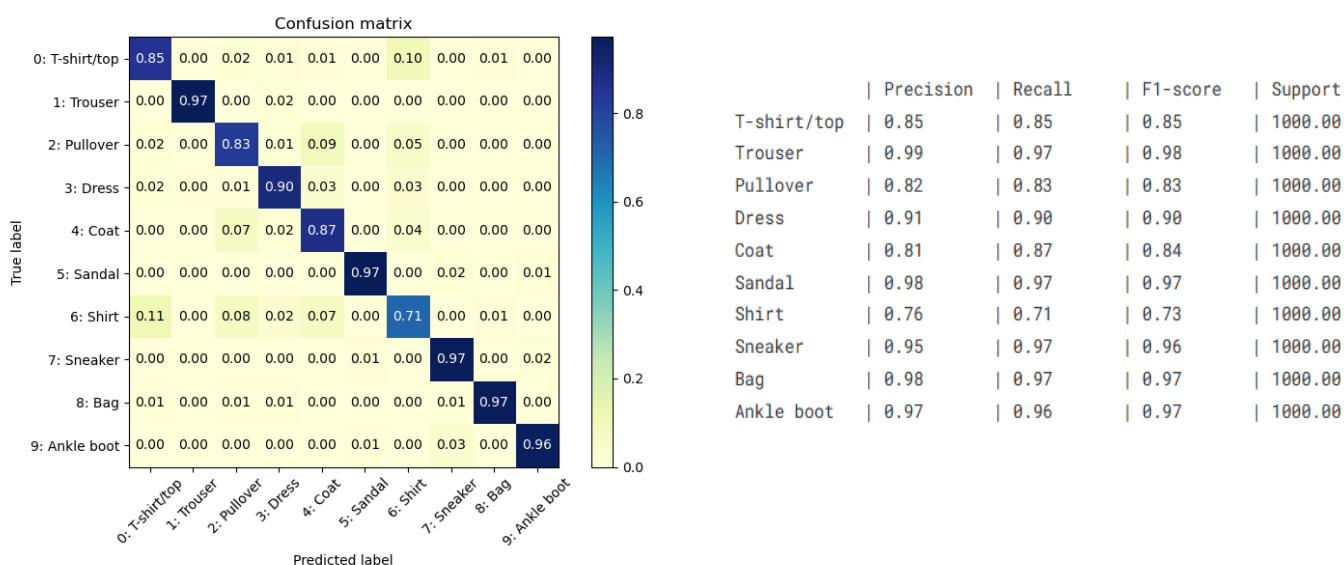
Below we show the confusion matrix and the full classification report per item:



# Multilayer Perceptron (MLP)

8

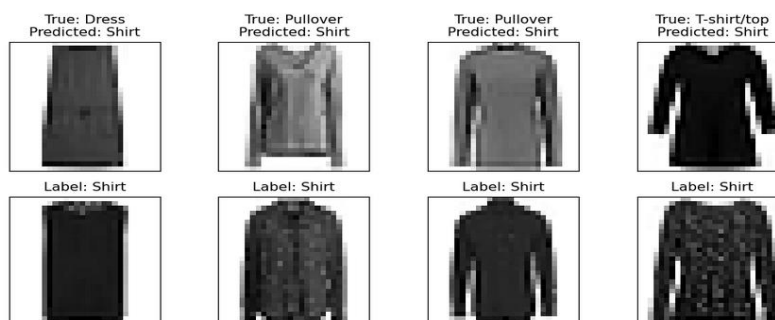
## Results' evaluation



Our intuition was right, and we see that some specific items are difficult to be matched correctly. This is why we proceeded in analyzing the incorrect classifications.

As seen from the confusion matrix, the errors happen most often on specific categories. We will explore these classes to get a better understanding of the behaviour of our model.

We created a plot containing 4 pictures of incorrect classification and 4 pictures of clothes corresponding to the predicted label. This will help us get a better understanding on the errors.



As seen from the confusion matrix, the Shirt category has the lowest recall and it is confused with the category T-Shirt/top, as well as the Pullover/Dress and coat. Looking at the plot above, we can see that this is expected, as many items of the top row (incorrect classifications) are similar to those of the second (predicted label pictures).

# Multilayer Perceptron (MLP)

9

Similarly, the T-shirt/top category is often confused with the Shirt and Pullovers are confused with Coats and Shirts, which can be expected, as seen from the images below.



## Convolutional Neural Network full implementation

In this part of the assignment we evaluated the performance of a CNN architecture.

- We first built a base model with 3 conv\_layers, which were used as a baseline
- Based on the results we got we created a sweep configuration to optimize hyperparameters

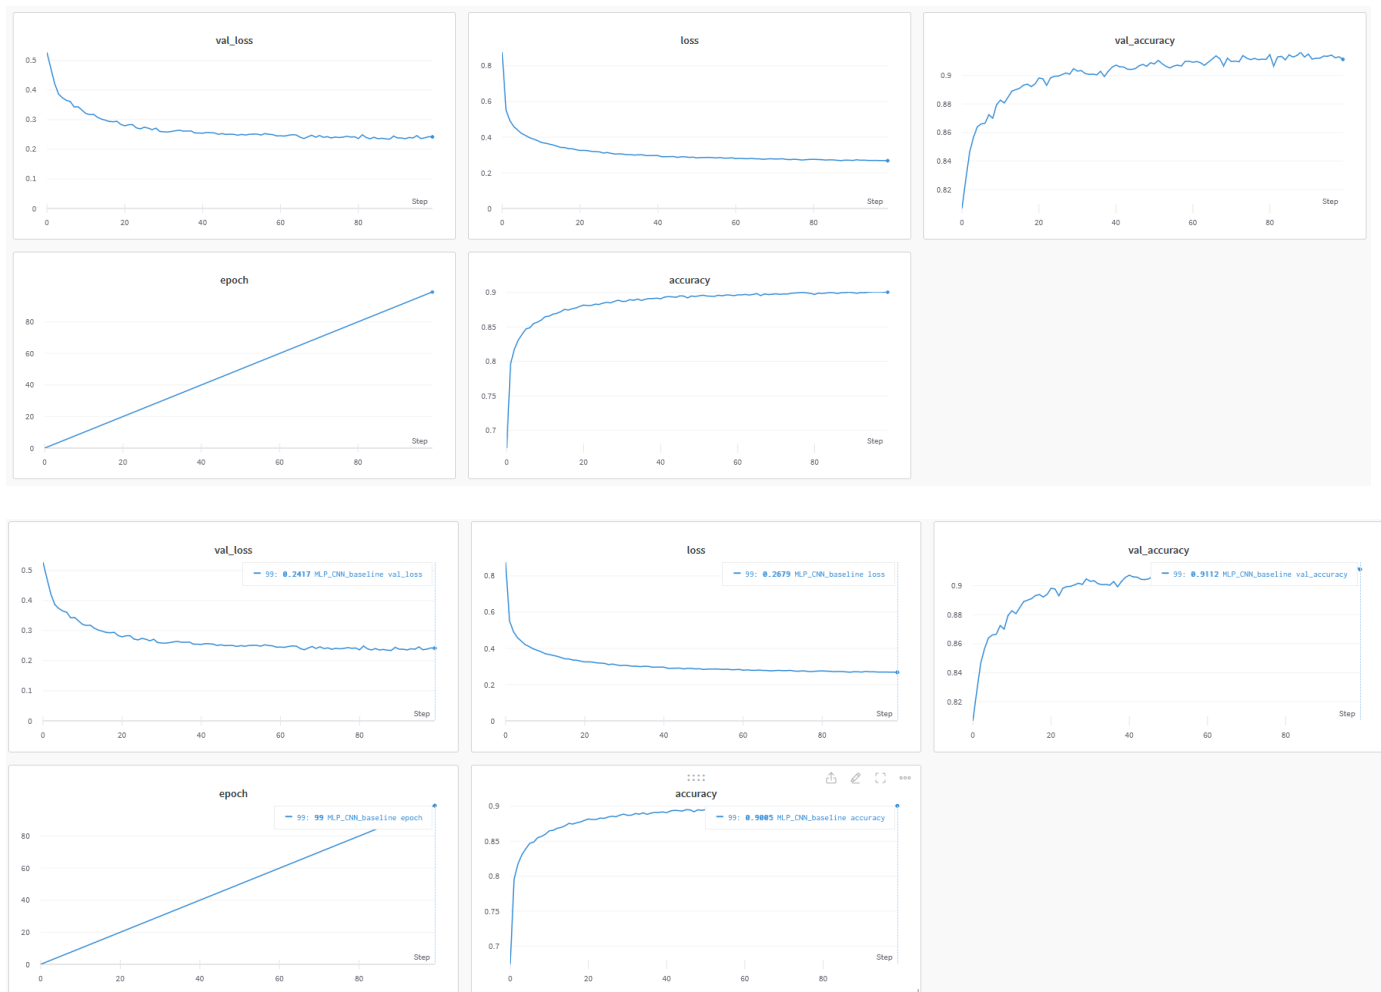
The base model with 3 conv\_layers:

```
for i in range(conv_layers):
    x = Conv2D(
        filters=8*(2**i),
        kernel_size=(3, 3),
        strides=(1, 1),
        padding='same',
        dilation_rate=(1, 1),
        activation=conv_activation,
        name='Conv2D-{:d}'.format(i + 1)
    )(x)
    x = MaxPool2D(
        pool_size=(2, 2),
        strides=(2, 2),
        padding='same',
        name='MaxPool2D-{:d}'.format(i + 1)
    )(x)
    if dropout:
        x = Dropout(
            rate=0.2,
            name='Dropout-{:d}'.format(i + 1)
        )(x)
# Flatten the convolved images so as to input them to a Dense Layer
x = Flatten(name='Flatten')(x)
```

# Convolutional Neural Network (CNN)

11

The results for Train Loss, Validation Loss, Test Loss, Train Accuracy, Validation Accuracy and Test Accuracy can be visualized effectively using once again Weights and Biases:




As we can see above, the model works surprisingly well, compared to the MLP one. We managed to match the accuracy of the fully optimized MLP model, with a relatively simple architecture. Nevertheless, we moved forward on fully optimizing the model.

We can also see, that the model reaches a plateau, without overfitting. This indicates that we should slightly increase the complexity of the model. Moreover, curves are not very smooth, thus we tried to decrease the learning rate. We built a new, more complex function, in order to create a new sweep which can be found in more detail in our code.

# Convolutional Neural Network (CNN)

12

We can also see our baseline CNN model compared with the MLP baseline models:



We can also see our baseline CNN model compared with the MLP baseline models: For the implementation of the more complex model, mentioned before, we implemented a hyper parameter tuning with the help of Weights and Biases once more.

We built a function named `train_model` which can build a model with the following blocks:

- Conv2D
- Batch Normalization
- Dropout
- 2D Max Polling

Each layer was built based on a list input, and depending on the value the Batch Normalization, Dropout and max polling layers could be skipped.

We continued with the step of generating random combinations of inputs. For each parameter we created a pool of inputs and generated random lists (of length 4) which were used as an input in the Weights and Biases sweep configuration.


Below you can see the the lists of parameters created for the configuration:

```
'parameters': {
  'batch_size': {'values': [64,
                           128,
                           512]},
  'filter_size': {'values': [[64, 256, 32, 16],
                             [16, 128, 1024, 16],
                             [128, 32, 128, 512],
                             [256, 512, 512, 512],
                             [256, 32, 64, 1024],
                             [1024, 512, 512, 128],
                             [256, 64, 128, 512],
                             [512, 512, 128, 64],
                             [64, 16, 512, 512],
                             [256, 32, 32, 128]]},
  'kernel_size': {'values': [[5, 3, 7, 5],
                              [3, 3, 3, 5],
                              [7, 7, 3, 5],
                              [3, 7, 5, 5],
                              [5, 5, 7, 5],
                              [5, 3, 7, 3],
                              [7, 5, 3, 7],
                              [7, 5, 5, 3],
                              [5, 5, 3, 5],
                              [3, 3, 7, 3]]},
  'batch_norm': {'values': [[True, True, False, False],
                             [False, True, False, True],
                             [False, True, True, True],
                             [False, True, False, False],
                             [False, False, True, True],
                             [True, False, False, False],
                             [False, False, False, False],
                             [True, False, True, False],
                             [True, True, True, False],
                             [False, False, True, False]]},
  'max_pooling_size': {'values': [[0, 0, 4, 2],
                                  [0, 4, 2, 4],
                                  [0, 0, 0, 4],
                                  [4, 0, 4, 0],
                                  [0, 2, 4, 0],
                                  [0, 4, 2, 2],
                                  [4, 0, 2, 0],
                                  [2, 2, 2, 2],
                                  [2, 4, 0, 4],
                                  [2, 0, 2, 0]]},
  'dropout': {'values': [[0.4, 0.0, 0.6, 0.2],
                          [0.4, 0.0, 0.0, 0.4],
                          [0.2, 0.2, 0.2, 0.6],
                          [0.0, 0.0, 0.0, 0.6],
                          [0.0, 0.6, 0.2, 0.4],
                          [0.0, 0.0, 0.2, 0.4],
                          [0.6, 0.4, 0.2, 0.2],
                          [0.4, 0.6, 0.6, 0.0],
                          [0.6, 0.4, 0.2, 0.0],
                          [0.0, 0.2, 0.6, 0.2]]},
  'learning_rate': {'values': [0.0001, 0.0005]}
```

# Convolutional Neural Network (CNN)

14

Having performed a large number of runs, we stopped further exploration, since the results did not seem to improve dramatically. Below you see the best parameter configuration for our model:



And running the predictions on the test data:

Model: "model"

Layer (type)	Output Shape	Param #
Input (InputLayer)	(None, 28, 28, 1)	0
Conv2D-1 (Conv2D)	(None, 28, 28, 512)	13312
MaxPool2D-1 (MaxPooling2D)	(None, 14, 14, 512)	0
Conv2D-2 (Conv2D)	(None, 14, 14, 512)	2359808
Conv2D-3 (Conv2D)	(None, 14, 14, 128)	3211392
BatchNormalization-3 (Batch Normalization)	(None, 14, 14, 128)	512
MaxPool2D-3 (MaxPooling2D)	(None, 7, 7, 128)	0
Conv2D-4 (Conv2D)	(None, 7, 7, 64)	204864
Dropout-4 (Dropout)	(None, 7, 7, 64)	0
Flatten (Flatten)	(None, 3136)	0
Output (Dense)	(None, 10)	31370

=====  
Total params: 5,821,258

Trainable params: 5,821,002

Non-trainable params: 256

None

313/313 [=====] - 14s 12ms/step - loss: 0.1921 - accuracy: 0.9361

---

Test Loss: 0.19210

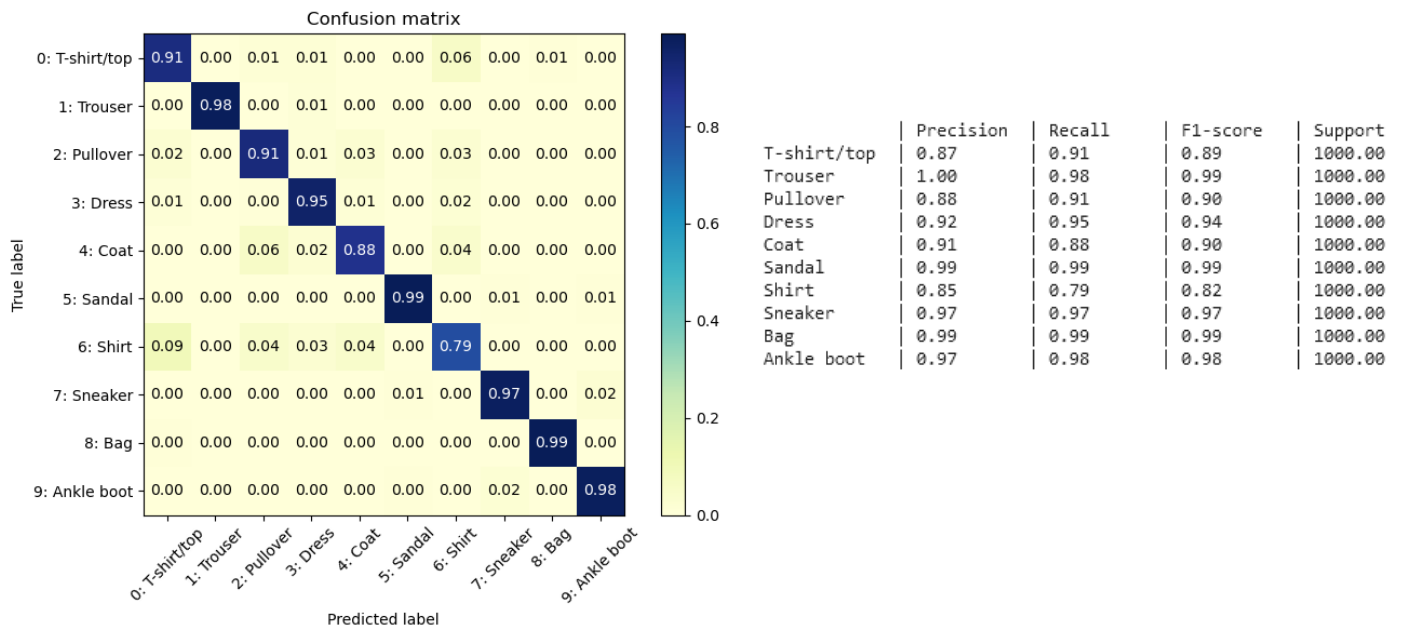
---

Test Accuracy: 0.93610

# Convolutional Neural Network (CNN)

15

Below we see the confusion matrix of our CNN and the full classification report per item:



Even though our model works pretty well we see (as in our MLP implementation), that some items probably are not correctly recognized by the CNN.

So below we see the incorrectly classified images in order to confirm the difficulties our model faced.






# Convolutional Neural Network (CNN)

16


True: Pullover  
Predicted: T-shirt/top




Label: T-shirt/top




True: Bag  
Predicted: T-shirt/top



Label: T-shirt/top




True: Bag  
Predicted: T-shirt/top




Label: T-shirt/top




True: Bag  
Predicted: T-shirt/top



Label: T-shirt/top




True: Shirt  
Predicted: Pullover




Label: Pullover




True: Coat  
Predicted: Pullover



Label: Pullover




True: Shirt  
Predicted: Pullover



Label: Pullover



True: Coat  
Predicted: Pullover



Label: Pullover

